

000 001 STARCRAFT II ARENA: EVALUATING LLMs IN 002 STRATEGIC PLANNING, REAL-TIME DECISION MAK- 003 ING, AND ADAPTABILITY 004 005

006 **Anonymous authors**
 007 Paper under double-blind review
 008
 009
 010

011 ABSTRACT 012

013 StarCraft II plays an important role in developing AI agents for real-time strategic
 014 reasoning due to its complex nature. However, people usually draw conclusions of
 015 how competent their agents are according to the level of the built-in agents in Star-
 016 Craft II which they can win in terms of the final success rate. Little intermediate
 017 quantitative information is considered while human-in-the-loop analysis is time
 018 inefficient, which results in inadequate reflection of the true strategic reasoning
 019 ability. In this work, we propose StarCraft II Arena, a well-designed benchmark
 020 for evaluating the strategic planning, real-time decision-making, and adaptability
 021 capabilities of large language models (LLMs) agents. We introduce using fine-
 022 grained capability metrics, allowing for targeted capture and analysis of specific
 023 capability, and further propose a detailed decision trace to enhance the under-
 024 standing of LLM behavior. We demonstrate the utility of such a benchmark by
 025 evaluating several state-of-the-art LLMs in various setups. Our results reveal dis-
 026 tinct performances in long-term strategy development, real-time decision-making,
 027 and adapting to environmental changes. Such results show that the StarCraft II
 028 Arena offers a deeper insight into the decision-making process of LLMs and has
 029 the potential to become a challenging and comprehensive benchmark for strategic
 030 reasoning.
 031

032 1 INTRODUCTION 033

034 LLMs have recently demonstrated exceptional capabilities in reasoning, planning, and problem-
 035 solving (Xi et al., 2023) across a range of domains, such as policy formulation (Xiao et al., 2023;
 036 Hua et al., 2023), investment decision-making (Weiss et al.; Li et al., 2023b), and strategic optimi-
 037 sation (Liu et al., 2024; Zhang et al., 2024a). Successfully completing these complex tasks requires
 038 intelligent agents to perceive, make decisions, and execute actions (Wooldridge & Jennings, 1995)
 039 within diverse and dynamic environments. This process not only involves deep reasoning to antic-
 040 ipate risks and weaknesses but also the ability to understand the motivations, beliefs, and potential
 041 deceptive behaviors of other agents (Hao et al., 2023; Premack & Woodruff, 1978; Street et al.,
 042 2024). Although LLMs have shown significant promise in managing such scenarios, positioning
 043 them as key technologies for achieving artificial general intelligence (AGI) (You et al., 2024; Morris
 044 et al.), their performance in real-world applications continues to face numerous challenges.

045 Evaluating the capabilities of LLM agents effectively is critical for the further development of this
 046 field. Traditional static evaluation datasets, while offering a standardized testing framework (Wang
 047 et al., 2019; Srivastava et al., 2022; Chen et al., 2021; Xie et al., 2024), are insufficient for captur-
 048 ing how models make decisions and adapt in dynamic environments. As a result, there has been
 049 growing interest in assessing the performance of large models within executable environments (Liu
 050 et al., 2023a; Xi et al., 2024)—simulated or real-world interactive platforms, including web naviga-
 051 tion (Lai et al., 2024), household tasks (Li et al., 2024), gaming (Bailis et al., 2024; Qi et al., 2024),
 052 and programming (Qian et al., 2024). Among these, games, with their clear rules and complex
 053 decision-making mechanisms, are considered ideal platforms for evaluating AI decision-making
 abilities (Costarelli et al., 2024).

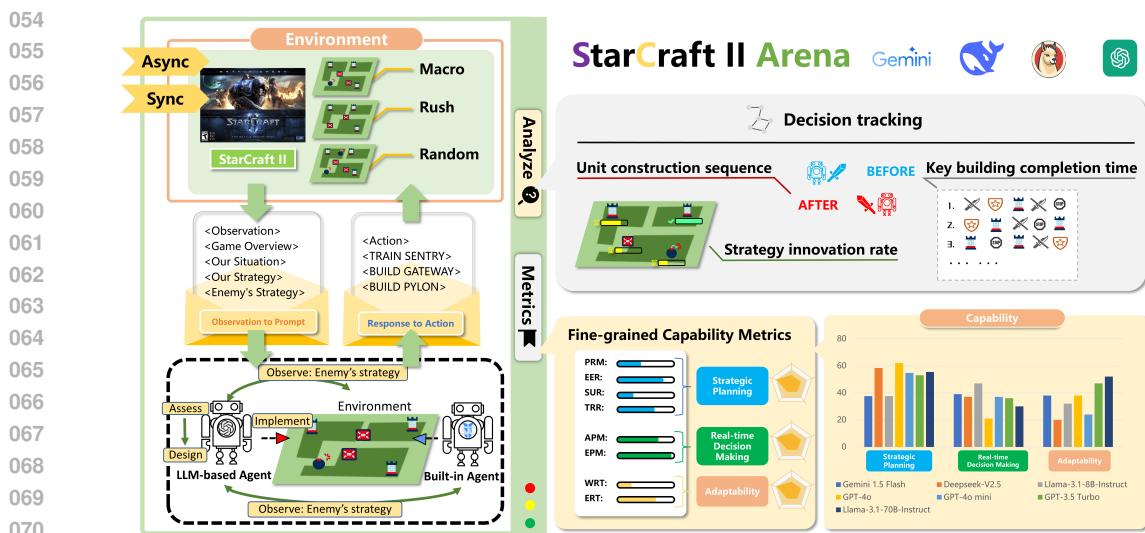


Figure 1: The overall framework of **StarCraft II Arena** which is designed to evaluate LLMs in strategic planning, real-time decision-making, and adaptability. It uses fine-grained capability metrics and a decision tracking system to capture key elements like Unit Construction Sequence and Strategy Innovation Rate, providing insights into LLMs’ decision-making and strategic reasoning.

However, existing evaluation benchmarks usually take the final outcome of a game such as success rate as the primary measure of performance (Duan et al., 2024). They further neglect the details of the intermediate outcomes gradually generated by LLM agents during the sequential decision-making process (Xi et al., 2024; Ma et al., 2024). Although the success rate reflects the overall ability of the competing agents, this singular metric is usually inadequate for reflecting how capable an LLM agent is when handling complexity or adapting to changing tasks. Therefore, a more comprehensive evaluation approach is required to better reflect the actual underlying reasoning capabilities of LLMs in dynamic environments.

In this study, we introduce **StarCraft II Arena**, a benchmark specifically designed to evaluate the abilities of LLMs for strategic planning, real-time decision-making, and adaptation in the game StarCraft II. As illustrated in Figure 1, unlike traditional benchmarks which usually depend on static tests or success rates, StarCraft II Arena incorporates fine-grained capability evaluation metrics, which allows for more detailed analysis of the performance of LLMs across multiple dimensions. Moreover, we also introduce a decision-tracking mechanism which records the intermediate decision results of LLMs during the task execution. By looking into the decision trajectories, we can analyze how LLMs adjust strategies in response to dynamic changes in the environment. As a result, it allows a more comprehensive understanding of the underlying decision process of LLMs rather than using the final outcome only. We demonstrate the utility of StarCraft II Arena by applying it to a range of recent LLM agents, both proprietary and open-source, and leading to the following key findings: (1) proprietary LLMs excel in long-term strategic planning and resource management but demonstrate limitations in dynamic environments that require rapid adaptation; (2) most existing LLMs struggle with handling incomplete information and adapting to rapidly evolving opponent strategies, limiting their ability to respond effectively to shifting tactics; and (3) smaller models show greater flexibility in real-time decision-making, particularly in high-frequency decision-making tasks, where they often outperform their larger counterparts. These findings highlight the potential of StarCraft II Arena as a challenging benchmark for LLMs agents in strategic reasoning tasks.

2 RELATED WORK

2.1 LLM-AS-AGENT

The application of large language models (LLMs) as agents is rapidly evolving, encompassing a diverse range of scenarios from single-agent to multi-agent systems (Xi et al., 2023). Early reinforcement learning (RL) agents learned through trial and error in complex environments, but they

108

109

110

111

112

113

114

115

116

Game	Imperfect Information	Strategic & Tactical	Dynamic space	Real-time v.s. Turn-based
Civilization(Wikipedia, 2024a)	✓	✗	✓	Turn-based
Dota 2(Wikipedia, 2024f)	✓	✗	✓	Real-time
Honor of Kings(Wikipedia, 2024c)	✓	✗	✓	Real-time
Diplomacy(Wikipedia, 2024b)	✗	✗	✗	Turn-based
Werewolf(Wikipedia (2024d))	✗	✗	✗	Turn-based
StarCraft II(Wikipedia, 2024e)	✓	✓	✓	Real-time

117

118

Table 1: Compare several games as LLM benchmarking environments based on four key dimensions: Imperfect Information, Strategic & Tactical, Dynamic Space, and Real-time v.s. Turn-based. Games like StarCraft II and Dota 2 feature imperfect information and dynamic spaces, which present significant challenges to the decision-making capabilities of LLMs. StarCraft II uniquely integrates both strategic and tactical elements, making it particularly suitable as a benchmark for assessing LLMs’ planning and decision-making abilities. Turn-based games like Civilization and Diplomacy provide a more controlled environment for long-term strategic planning, while real-time games impose time constraints that test the models’ ability to react swiftly.

119

120

121

122

123

124

125

126

127

were typically suited only for highly structured tasks and required substantial training time and data (Pourchot & Sigaud, 2018). In contrast, LLM-based agents, trained on extensive text datasets, possess strong language understanding, instruction-following, and generation capabilities (Liu et al., 2022; Lu et al., 2023), enabling them to flexibly navigate varied situations and demonstrate few-shot and zero-shot generalization abilities (Wei et al., 2021; Yao et al., 2022), thus achieving seamless task transfer. Furthermore, these LLM agents exhibit advanced cognitive abilities akin to human intelligence, including chain-of-thought reasoning (Wei et al., 2022; Jin & Lu, 2023; Zhang et al., 2023), planning (Huang et al., 2024a), self-reflection (Madaan et al., 2024), memory (Zheng et al., 2023a; Zhang et al., 2024b), and learning (Zhang et al., 2024a; Xi et al., 2024). These capabilities empower LLM agents to effectively tackle complex decision-making scenarios.

128

129

130

131

132

133

134

135

136

137

138

139

140

141

142

143

In multi-agent systems, LLMs must not only interact with their environment but also engage in effective communication and collaboration among multiple agents to accomplish tasks (Pourchot & Sigaud, 2018). Such systems emphasize the importance of agent communication and cooperation, allowing them to operate within dynamic and complex environments, such as game simulations (Xu et al., 2023), financial market analysis (Chen et al., 2023), and software development (Qian et al., 2024). Strategic reasoning is particularly crucial in this context, as it requires agents to understand and predict the actions of other agents and adjust their strategies accordingly.

144

145

2.2 BENCHMARKS FOR AI AGENTS

146

147

148

149

150

151

152

153

154

155

156

157

158

159

160

161

Evaluation Environments. In previous research, the capabilities of large models have primarily been assessed through the construction of static datasets (Wang et al., 2019; Srivastava et al., 2023; Zheng et al., 2023b; Yue et al., 2024; Xie et al., 2024). While an increasing number of benchmarks have introduced broader tasks and datasets, most remain confined to traditional tasks and fail to comprehensively evaluate the capabilities of large language models (LLMs) in open-ended generation, multi-turn interactions, and agent-based roles (Gur et al., 2023; Huang et al., 2024b; Liu et al., 2023b). As LLMs become more adept at addressing real-world challenges, there is a growing trend towards evaluation methods that are based on executable environments rather than static datasets (Gur et al., 2023; Wang et al., 2023; Shinn et al., 2024). Specifically, researchers are now focusing on areas such as web navigation (Deng et al., 2024; Yao et al., preprint), text-based games (Bailis et al., 2024; Mukobi et al., 2023), household tasks (Wang et al., 2022), digital games (Qi et al., 2024; Ma et al., 2023), avatar tasks (Han et al., 2024), tool usage (Tang et al., 2023), and programming (Qian et al., 2024; Zheng et al., 2023b), all of which provide a more realistic context for assessing LLMs. In particular, games are widely regarded as ideal experimental platforms for evaluating the decision-making capabilities of large models (Liu et al., 2023a). By placing models in dynamic and complex gaming environments, researchers can effectively gauge their performance in real-world scenarios.

With clear rules and flexible customization, **gaming environments** have been applied widely for evaluating AI decision-making abilities. We compare several popular games as LLM benchmarking environments in Table 1 based on four key dimensions: Imperfect Information, Strategic & Tactical, Dynamic Space, and Real-time v.s. Turn-based. Notably, StarCraft II, as a complex real-time strategy game (Vinyals et al., 2019; Samvelyan et al., 2019), provides an ideal platform for evaluating LLMs’ capabilities in strategic reasoning and multi-agent interaction. By assessing LLM performance in this environment, we can gain deeper insights into how these models respond to complex decision-making and dynamic changes.

Evaluation Metrics. Some studies employ game-theoretic tools to systematically evaluate the decision-making abilities of large models within games, aiming to measure their strategic choices and adaptability (Duan et al., 2024). However, these studies often concentrate on simple games with a single dimension, failing to fully capture the complexity of the models’ decision-making processes. Additionally, other research has focused on dissecting the capabilities of large models to explore performance variations and potential advantages across different gaming scenarios (Wu et al., 2023; Ma et al., 2024). Nevertheless, most analyses predominantly emphasize win rates, lacking fine-grained capability metrics and decision trajectory analyses, which limits a comprehensive understanding of the models’ performance (Costarelli et al., 2024; Duan et al., 2024; Liu et al., 2023a; Wu et al., 2023).

3 PRELIMINARY

The agent’s interaction with the environment in StarCraft II is modeled as a Partially Observable Markov Decision Process (POMDP), defined by the tuple $\langle W, S, A, O, T \rangle$, where W represents the victory goal, S is the state space, A is the valid actions space, O is the observation space (including environmental feedback), and T is the state transition function. The agent interacts with the environment by selecting actions from A based on the current state S and observations O , with the state evolving according to T .

Two-level inference. In StarCraft II, the complexity of reasoning arises from the need to handle a large observation space and multi-dimensional strategic tasks. This requires two levels of reasoning: high-level strategic planning, such as resource management and army mobilization, and low-level decision-making, such as micro-control in local battles.

$$p_{\pi}(\tau) = p(s_0) \prod_{t=0}^{T-1} p(a_t^{high}|s_t, c^{high}) \cdot p(a_t^{low}|s_t, a_t^{high}, c^{low}) \cdot T(s_{t+1}|s_t, a_t^{low}, f_t) \quad (1)$$

Here, a_t^{high} represents a high-level decision based on the global strategy c^{high} , and a_t^{low} is a low-level action based on local feedback c^{low} . The state transition function $T(s_{t+1}|s_t, a_t^{low}, f_t)$ models how the environment transitions in response to the agent’s low-level actions and the feedback received from the environment.

3.1 FINE-GRAINED CAPABILITY METRICS

To evaluate the specific capabilities of the Large Language Model (LLM), each model is tested in m scenarios, and performance is assessed based on aggregated metrics. A final capability score is calculated as follows:

$$T = \sum_{i=1}^m W_i \cdot \beta_i \cdot \left(\frac{1}{n} \sum_{j=1}^n \frac{\bar{R}_{y_j} - \mu_j}{\sigma_j} \right) \quad (2)$$

Here, W_i represents the weight of scenario i , and β_i is a moderating factor for scenario i . For each metric j , \bar{R}_{y_j} is the average result across k_j runs, μ_j and σ_j are the mean and standard deviation of the metric, used for normalization. The final score is a weighted sum of normalized metrics across all scenarios.

Capacity	Metrics	Scene selection	
		Opponent strategy	Operation mode
Strategic Planning	RPM, EER, SUR, TRR	Macro	Async/Sync
Real-time Decision Making	APM, EPM	Rush	Async/Sync
Adaptability	WRT, ERT	Random	Async/Sync

Table 2: Outlines the capacity and metrics utilized in the StarCraft II Arena benchmark for evaluating large language models (LLMs). The table highlights three key dimensions: Strategic Planning, Real-time Decision-Making, and Adaptability. Each dimension is associated with specific metrics, such as Resource Management Ability (RMA), Resource Utilization Efficiency (RUE), Actions Per Minute (APM), and Win Rate Growth Rate (WRGR), among others. Additionally, the table details the scene selection strategies, including opponent strategies like Macro, Rush, and Random, along with the operational modes categorized as Async or Sync. This comprehensive structure facilitates a detailed assessment of LLM capabilities within complex strategic environments.

4 STARCRAFT II ARENA - OVERVIEW

StarCraft II Arena is a benchmark specifically designed to assess the performance of various LLMs in the strategic real-time game Starcraft II. It evaluates the capabilities of LLMs from the perspectives of strategic planning, real-time decision-making, and adaptability through a series of carefully constructed gaming scenarios. In contrast to traditional evaluation methods, StarCraft II Arena offers more refined quantitative analysis metrics and an additional behavior-tracking mechanism, allowing for a deeper, multi-faceted understanding of the underlying reasoning process of LLMs. We shall explain the detailed capability dimensions, the fine-grained capability metrics, the design of different testing scenarios, and the behavioral analysis using decision tracking as follows. Table 2

4.1 DECOMPOSITION OF THE STRATEGIC REASONING CAPABILITY

LLM agents demonstrate several advanced cognitive abilities akin to human intelligence in complex environments, including chain-of-thought reasoning(Wei et al., 2022; Jin & Lu, 2023; Zhang et al., 2023), planning(Huang et al., 2024a), self-reflection(Madaan et al., 2024), memory(Zheng et al., 2023a; Zhang et al., 2024b), and learning(Zhang et al., 2024a; Xi et al., 2024). Based on these characteristics, the selection of strategic planning, real-time decision-making, and adaptability as core evaluation dimensions is logically grounded. These three dimensions encapsulate the essential capabilities required for agents to tackle complex tasks, representing holistic thinking, rapid response, and flexible adaptation.

Strategic planning serves as the foundation for LLMs when addressing long-term objectives in dynamic environments. Short-term reactions alone are insufficient to manage fluctuating conditions. The model must have a broad view, ensuring the efficient allocation of resources, the prioritisation of tasks, and the formulation of long-term strategies to maintain a competitive edge. Effective strategic planning demands not only the ability to foresee potential future developments but also to make informed decisions concerning resource management, technological advancements, and unit production, thereby securing and sustaining a strategic advantage. In games like StarCraft II, for instance, an LLM must efficiently manage early resource accumulation and expansion while preparing for large-scale combat in the mid to late game.

Real-time decision-making is critical when using LLMs in real applications. While several LLM agents claim to be capable for complex tasks, some are evaluated in a setup where the executing testing system needs to be suspended while the LLM agents perform inference. A perfect strategic plan would be ineffective if the model takes too long to perform inference and cannot respond rapidly enough to the changing conditions. The model must continuously process dynamic information and adjust its tactics accordingly. For instance, during a sudden enemy assault, the LLM must promptly deploy units to defend or counter-attack, maintaining control over the situation. This capability requires not only rapid information processing but also the ability to evaluate multiple strategies quickly and efficiently to preserve overarching strategic objectives.

270 **Adaptability** determines the model’s ability to remain competitive in evolving environments. As
 271 opponent strategies, resource conditions, and task priorities change, an adaptive model can adjust its
 272 approach based on previous feedback, refining its strategies to address new challenges. This reflects
 273 the model’s flexibility and its ability to learn from experience. Quantitative metrics like win-rate
 274 growth and error-rate reduction measure how well the model improves its decision-making over
 275 time, ensuring a sustained advantage in long-term gameplay.

277 4.2 EVALUATION METRICS FOR THE INDIVIDUAL CAPABILITIES

279 Recent studies have highlighted that using success rate as the primary metric for agent evaluation
 280 fails to capture the nuanced differences in how language model agents perform partial tasks (Liu
 281 et al., 2023a; Li et al., 2023a). In adversarial games such as StarCraft II, this approach does not dis-
 282 tinguish between the success of local tactics and the failure to achieve overall victory, instead treating
 283 all instances of not reaching the final objective as failures. This overlooks the agent’s incremental
 284 achievements or the effectiveness of its local strategies during gameplay. Although alternative met-
 285 rics like reward scores can be used to assess performance, the lack of standardisation complicates
 286 cross-environment comparisons, limiting their broader applicability(Chevalier-Boisvert et al., 2018;
 287 Wang et al., 2022; Hausknecht et al., 2020).

288 To address these issues, we introduce Fine-Grained Capability Metrics to provide a more precise
 289 evaluation of LLM performance across different task stages. Quantitative metrics are used to eval-
 290 uate specific competencies such as resource management, real-time decision-making, and adapt-
 291 ability. For instance, the Resource Management Ability (RPM) is calculated by summing the total
 292 minerals and vespene gas collected during the game, reflecting the model’s efficiency in resource
 293 gathering. The formula is:

$$294 \quad RPM_i = \sum_{t=1}^T (collected_minerals_i(t) + collected_vespene_i(t)) \quad (3)$$

298 Similarly, the Supply Utilization Rate (SUR) measures the ratio of supply used to maximum supply
 299 capacity, offering insight into the model’s ability to effectively produce units. The formula is:

$$300 \quad SUR_i = \frac{\sum_{t=1}^T supply_used_i(t)}{\sum_{t=1}^T supply_cap_i(t)} \quad (4)$$

303 These metrics capture the model’s performance across resource allocation, unit production, and tech-
 304 nological development, providing a quantitative basis for evaluating strategic planning capabilities.
 305 A full list of these quantitative metrics, along with their respective formulas, will be provided in the
 306 appendix for reference.

308 4.3 DECISION TRACKING AND BEHAVIORAL ANALYSIS

310 To gain a deeper understanding of the decision-making processes and behavioural patterns of LLMs
 311 in the StarCraft II Arena, this paper introduces a decision tracking and behavioural analysis sys-
 312 tem. This system records and analyzes critical operations and decisions made by the LLM during
 313 gameplay, contextualizing them within the current game state and mission objectives to uncover the
 314 underlying strategic logic.

315 Specifically, the system captures key actions such as resource allocation, technological upgrades,
 316 unit production, and troop movements. A decision trace of an LLM consists of a chronological
 317 record of these key actions and decisions, detailing the context in which they were made. It typically
 318 comprises three main components: (1) Action Type, indicating the specific type of decision made
 319 (e.g., resource allocation, unit production); (2) Decision Context, which includes the game state and
 320 mission objectives at the time of the decision; and (3) Outcome, reflecting the consequences of the
 321 decision on the gameplay.

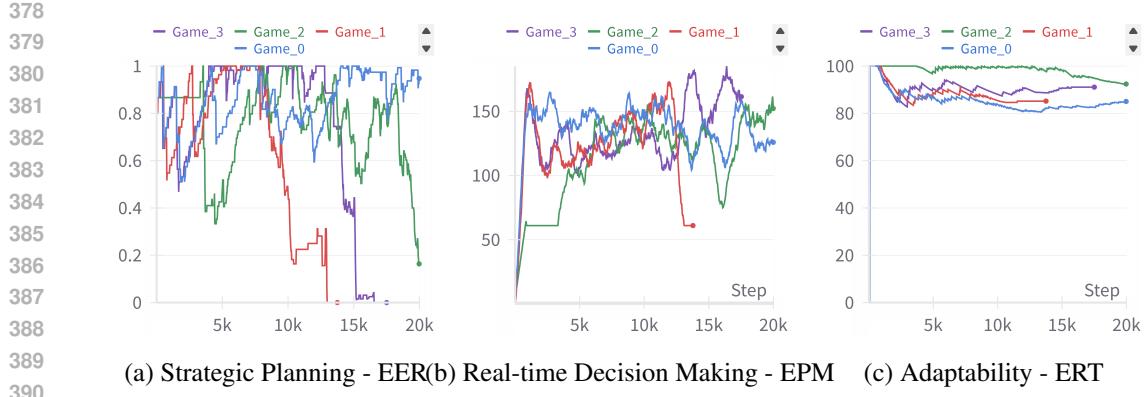
322 For example, a decision trace might document an LLM’s choice to produce a specific unit type in
 323 response to an opponent’s strategy, including the unit produced, the rationale based on the game
 324 state, and the subsequent impact on the match outcome. Screenshots illustrating these decision

Stage	Trajectory
324	
325	
326	
327	
328	
Early	<p>Decision: 0: <TRAIN PROBE>→1: <BUILD PYLON>→2: <BUILD NEXUS>→3: <TRAIN PROBE>→4: <BUILD ASSIMILATOR> Our Strategy: Expand the number of workers, build additional resource collection structures, and ensure base defense. Suggestions: Focus on building additional Probes to boost resource gathering, construct more Pylons to increase supply cap, and consider expanding to new resource locations. Step:429</p> <p>Decision: 0: <TRAIN STALKER>→1: <BUILD GATEWAY>→2: <TRAIN PROBE>→3: <BUILD ASSIMILATOR>→4: <RESEARCH PROTOSSGROUNDARMORSLEVEL1> Our Strategy: Build a stable economy and gradually advance technology and the army. Suggestions: Continue to build up the economy while gradually increasing army supply and advancing technology to prepare for mid-game engagements. Step:4629</p> <p>Decision: 0: <BUILD GATEWAY>→1: <BUILD SHIELDBATTERY>→2: <TRAIN SENTRY>→3: <RESEARCH PROTOSSGROUNDARMORSLEVEL2>→4: <SCOUTING PROBE> Our Strategy: Build a solid economic foundation, maintain a basic military presence, and invest in early upgrades to enhance force capabilities. Suggestions: Consider diversifying our army composition by introducing more advanced units like Sentries or Adepts to complement the existing Zealots and Stalkers. Focus on scouting to gather intelligence on the enemy's strategy. Step:7004</p>
329	
330	
331	
332	
333	
334	
Mid	<p>Decision: 0: <TRAIN SENTRY>→1: <BUILD ROBOTICSFACILITY>→2: <TRAIN IMMORTAL>→3: <BUILD ASSIMILATOR>→4: <RESEARCH PROTOSSGROUNDARMORSLEVEL2> Our Strategy: Improve the effectiveness of the army around balanced economy, unit composition and early upgrades. Suggestions: Consider diversifying your unit composition to include more advanced units like Sentries or Immortals to counter potential enemy threats. Additionally, continue upgrading your technology to stay ahead in battles. Step:8454</p> <p>Decision: 0: <TRAIN IMMORTAL>→1: <TRAIN OBSERVER>→2: <TRAIN VOIDRAY>→3: <BUILD GATEWAY>→4: <BUILD ASSIMILATOR> Our Strategy: Focus on a mix of ground and air units, supported by upgrades. We are also investing in infrastructure for advanced unit production. Suggestions: Consider diversifying our army composition by training units from the Robotics Facility and Stargates to counter a wider range of enemy units. Step:9554</p> <p>Decision: 0: <TRAIN HIGHTEMPLAR>→1: <BUILD ROBOTICSBAY>→2: <RESEARCH EXTENDEDTHERMALLANCE>→3: <BUILD ASSIMILATOR>→4: <TRAIN OBSERVER> Our Strategy: Focus on upgrading units and technology while maintaining a balanced economy and military composition. Suggestions: Consider expanding unit diversity by incorporating more advanced units like Colossi or High Templars to enhance our army's capabilities. Additionally, focus on tech upgrades to stay ahead in battles. Step:10529</p>
335	
336	
337	
338	
339	
340	
341	
342	
343	
Mid to Late	<p>Decision: 0: <BUILD GATEWAY>→1: <BUILD PHOTONCANNON>→2: <TRAIN STALKER>→3: <RESEARCH PROTOSSAIRARMORSLEVEL2>→4: <EXPAND TO NEW RESOURCE LOCATION> Our Strategy: Around the mix of gateway units and technical units such as Stalkers, Immortals, and Voidray. Suggestions: Enhance our unit composition by adding more production structures like Gateways to increase unit production capacity. Consider expanding tech options to counter potential enemy strategies effectively. Step:14129</p> <p>Decision: 0: <TRAIN STALKER>→1: <BUILD GATEWAY>→2: <BUILD SHIELDBATTERY>→3: <RESEARCH PROTOSSGROUNDARMORSLEVEL2>→4: <SCOUTING OBSERVER> Our Strategy: Around the unit diversity and balance method of technological progress. The emphasis on stalkers suggests a flexible military composition capable of handling a variety of threats. Suggestions: Consider diversifying our unit composition further to adapt to potential enemy strategies. Focus on maintaining map control and scouting to stay ahead of the opponent. Step:17554</p> <p>Decision: 0: <TRAIN STALKER>→1: <BUILD GATEWAY>→2: <CHRONOBOOST CYBERNETICSCORE>→3: <RESEARCH PROTOSSAIRWEAPONS-LEVEL2>→4: <TRAIN IMMORTAL> Our Strategy: Focus on balanced unit composition and technology upgrades to improve the effectiveness of our forces. Suggestions: Consider diversifying our army composition further to counter the enemy's Stalkers effectively. Focus on unit production and upgrades to strengthen our army. Step:20929</p>
344	
345	
346	
347	
348	
349	

Table 3: Presents a detailed decision trajectory analysis of an LLM’s gameplay in StarCraft II across various stages: Early, Mid, and Mid to Late. The table outlines specific decisions made by the model, including the training of units, building structures, and conducting research. Each entry details the decision-making process, the associated strategy, and suggestions for optimizing performance. By capturing these trajectories, the table illustrates how the LLM navigates complex strategic choices, adapts to the game environment, and develops its military and economic strategies over time, providing insights into its strategic reasoning capabilities.

traces will be provided in the appendix to offer visual clarity. These actions are traced throughout the entire decision chain to assess the coherence of the LLM’s strategic planning, flexibility in tactical adjustments, and adaptability to opponent strategies. Additionally, the system employs visualization tools to present real-time behavioural pathways, allowing researchers to observe how the model reacts in various scenarios. This real-time tracking provides insights into effective decision-making patterns and potential areas for optimization, offering a comprehensive understanding of both tactical execution and strategic intent.

In parallel, qualitative metrics are introduced to further analyze the LLM’s decision-making patterns and strategic innovations. These metrics are closely tied to the decision tracking system, capturing specific game actions such as the Unit Construction Order, which reflects the model’s tactical priorities by documenting the sequence in which units are built. For example, prioritizing basic units may indicate a focus on early offensive strategies, whereas building high-tech units suggests a defensive or late-game approach. Similarly, the Key Building Completion Time records the timing of critical structures, such as Nexus or Gateway, to determine whether the construction order supports economic growth or military objectives. Moreover, the Strategy Innovation Rate, which measures the frequency of adopting new strategies across multiple games, is derived from continuous monitoring of strategic shifts. These qualitative metrics provide a detailed view of the model’s adaptability and capacity for innovation across various game scenarios. Definitions and methods for these metrics will be further detailed in the figure 4.3, along with examples and screenshots to effectively illustrate the decision tracking process.



(a) Strategic Planning - EER (b) Real-time Decision Making - EPM (c) Adaptability - ERT

Figure 2: Performance indicators for evaluating LLM capabilities in StarCraft II: (a) Strategic Planning - EER (Efficiency of Resource Utilization), (b) Real-time Decision Making - EPM (Effective Actions Per Minute), and (c) Adaptability - ERT (Error Rate Trend). Each graph displays the performance trends of different game sessions (Game_0 , Game_1 , Game_2 , Game_3) over time steps.

5 ARENA EVALUATION

5.1 EVALUATION SETUP

We conducted a comprehensive evaluation of popular large language models, including both proprietary API-based models and open heavyweight models. Firstly, we report the success rates and progress rates of these agents. Subsequently, we provide a detailed analysis of their performance and measure the various capabilities of the LLM agents, culminating in a further analysis of their decision-making trajectories.

5.2 MAIN RESULTS

Model	Win Rate	Strategic Planning	Real-Time Decision	Adaptability	Overall Score
GPT-4o(OpenAI, 2024b)	2/10	62.01	21.12	38.64	57758
GPT-4o mini(OpenAI, 2024a)	5/10	54.71	37.51	24.52	62541
GPT-3.5 Turbo(OpenAI, 2023)	4/10	53.24	36.23	47.41	60914
Gemini 1.5 Flash(Reid et al., 2024)	5/10	37.56	39.34	38.18	55940
DeepSeek-V2.5(DeepSeek-AI, 2024)	2/10	58.35	37.11	20.16	43070
Llama-3.1-8B-Instruct(Dubey et al., 2024)	3/10	37.56	47.05	32.71	44901
Llama-3.1-70B-Instruct(Dubey et al., 2024)	2/10	55.44	30.24	52.77	46825

Table 4: Demonstrates the performance of several large-scale language models on different ability dimensions, specifically win rate, strategic planning, social reasoning, real-time decision making, teamwork, learning ability, and overall score.

Fine-grained capability metrics provide a more detailed and insightful evaluation of model performance than simple success rates. These metrics reveal substantial differences in how models handle strategic planning, real-time decision-making, and adaptability. GPT-4o achieved the highest score in strategic planning with 62.01 points, showcasing its strength in long-term resource management and strategy. However, its real-time decision-making score of 21.12 points was notably lower, indicating slower response times to in-game events. Conversely, Llama 3.1 Instruct 8B excelled in real-time decision-making with a score of 47.05, yet its strategic planning score was lower at 37.56, suggesting it is better suited to making quick decisions under pressure rather than managing long-term strategies. Llama 3.1 Instruct 70B led in adaptability, particularly in metrics like win rate growth and error rate reduction, which reflects its ability to learn and adjust to evolving game conditions. These findings demonstrate that fine-grained metrics enable a more nuanced understanding of each model’s strengths and weaknesses, beyond what win rates alone can offer.

Evaluating performance in both synchronous and asynchronous settings reveals how time constraints impact model behaviour. In synchronous settings, where rapid decision-making is essential, Llama 3.1 Instruct 8B and GPT-4o mini performed exceptionally well, with real-time decision-making scores of 47.05 and 37.51, respectively, highlighting their ability to respond quickly to changing conditions. However, in asynchronous settings, where models have more time to process information and plan their strategies, GPT-4o and DeepSeek-V2.5 excelled, achieving strategic planning scores of 62.01 and 58.35, respectively. This contrast illustrates that models adept at quick decision-making may face challenges in handling complex, long-term planning, while those strong in strategic planning may be slower to react in time-critical situations. Therefore, considering both settings is crucial for a comprehensive evaluation, as it underscores the balance between short-term reactivity and long-term planning in model performance.

Closed-source models consistently outperform open-source models in strategic planning and overall performance. GPT-4o and GPT-4o mini were the top performers, with GPT-4o achieving the highest overall score of 57,758, significantly surpassing open-source models such as Llama 3.1 Instruct 70B and Llama 3.1 Instruct 8B, which scored 46,825 and 44,901, respectively. This demonstrates that closed-source models benefit from larger training datasets and more optimised architectures, giving them an advantage in resource management and strategic tasks. Despite this, open-source models like Llama 3.1 Instruct 8B showed competitive performance in real-time decision-making, scoring 47.05, suggesting that open models are better suited for tasks requiring rapid responses. While closed-source models dominate in long-term planning and complex reasoning, the performance of open-source models in real-time decision-making highlights their potential, especially with further optimisation and development. This suggests that with additional resources, open-source models could narrow the performance gap, particularly in more complex strategic tasks.

5.3 ANALYTICAL EVALUATION

Unit Construction Order. The Unit Construction Order is a critical metric that reflects the tactical priorities of the LLM. As illustrated in Table 4.3, the decision trace reveals the sequence of units constructed during gameplay, allowing us to assess strategic intent. For instance, in Game 3, a notable shift occurs as the game progresses into the later stages, with the LLM beginning to prioritize the construction of advanced units. This transition can be observed through the timing of unit production, which indicates that the model is adapting its strategy in response to the evolving game dynamics. For example, the LLM initially focuses on building basic units, which is typical in the early game to establish a strong economy and military presence. However, as the game advances, there is a marked increase in the production of higher-tier units, such as Colossi and High Templars. This suggests a strategic shift aimed at countering opponent threats and enhancing combat effectiveness. The visual representation in Figure 2 further emphasizes this point, showing the timing and frequency of unit production across different game sessions.

Key Building Completion Time. The Key Building Completion Time metric assesses the efficiency and timing of critical structures necessary for advancing the game’s strategy. In Game 3, we observe that the completion of vital buildings such as the Robotics Facility and Templar Archives coincides with the shift towards producing more advanced units. This timing indicates that the LLM is effectively managing its resources to maximize its strategic output. For instance, if the completion time for these structures is relatively short and aligns with the LLM’s decision to construct advanced units, it reflects a well-coordinated strategy that prioritizes technological advancement alongside unit production. This synchronization is crucial for maintaining pressure on opponents and capitalizing on strategic opportunities.

Strategy Innovation Rate. The Strategy Innovation Rate measures the frequency with which the LLM adopts new strategies during gameplay. By analyzing the decision traces, we can identify instances where the model implements novel tactics or unit combinations in response to evolving game conditions. For example, in the later stages of Game 3, the LLM demonstrates an increase in strategic innovation, as evidenced by its willingness to experiment with unit compositions that differ from those used in earlier phases. This adaptability is highlighted in Figure 2, where we can see fluctuations in performance metrics over time. Such fluctuations suggest that the LLM is actively refining its strategies to better respond to opponents and the overall game state. This capacity for innovation is a testament to the model’s robust decision-making framework, enabling it to remain competitive in a dynamic environment.

486 6 CONCLUSION
 487

488 This study presents StarCraft II Arena as a comprehensive benchmark for assessing the capabilities
 489 of LLMs in strategic planning, real-time decision-making, and adaptability. The findings demon-
 490 strate that LLMs possess varying strengths across these dimensions, with notable performance in
 491 strategic reasoning and adaptability. By employing fine-grained metrics, we highlight the limita-
 492 tions of traditional success rates in capturing the true decision-making processes of LLMs. Our
 493 analysis reveals that models like GPT-4o excel in long-term strategic planning, while others, such as
 494 Llama 3.1 Instruct 8B, exhibit superior real-time decision-making capabilities. This detailed eval-
 495 uation not only enhances our understanding of LLMs' cognitive abilities in complex environments
 496 but also lays the groundwork for future advancements in AI research, emphasizing the importance
 497 of dynamic assessments in evaluating AI agents.

498 REFERENCES
 499

- 500 Suma Bailis, Jane Friedhoff, and Feiyang Chen. Werewolf arena: A case study in llm evaluation via
 501 social deduction. *arXiv preprint arXiv:2407.13943*, 2024.
 502
- 503 Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde De Oliveira Pinto, Jared
 504 Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large
 505 language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- 506 Yiting Chen, Tracy Xiao Liu, You Shan, and Songfa Zhong. The emergence of economic rationality
 507 of gpt. *Proceedings of the National Academy of Sciences*, 120(51):e2316205120, 2023.
 508
- 509 Maxime Chevalier-Boisvert, Dzmitry Bahdanau, Salem Lahlou, Lucas Willems, Chitwan Saharia,
 510 Thien Huu Nguyen, and Yoshua Bengio. Babyai: A platform to study the sample efficiency of
 511 grounded language learning. *arXiv preprint arXiv:1810.08272*, 2018.
- 512 Anthony Costarelli, Mat Allen, Roman Hauksson, Grace Sodunke, Suhas Hariharan, Carlson Cheng,
 513 Wenjie Li, and Arjun Yadav. Gamebench: Evaluating strategic reasoning abilities of llm agents.
 514 *arXiv preprint arXiv:2406.06613*, 2024.
- 515 DeepSeek-AI. Deepseek-v2: A strong, economical, and efficient mixture-of-experts language
 516 model, 2024.
 517
- 518 Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen, Sam Stevens, Boshi Wang, Huan Sun, and Yu Su.
 519 Mind2web: Towards a generalist agent for the web. *Advances in Neural Information Processing
 520 Systems*, 36, 2024.
- 521 Jinhao Duan, Renming Zhang, James Diffenderfer, Bhavya Kailkhura, Lichao Sun, Elias Stengel-
 522 Eskin, Mohit Bansal, Tianlong Chen, and Kaidi Xu. Gtbench: Uncovering the strategic reasoning
 523 limitations of llms via game-theoretic evaluations. *arXiv preprint arXiv:2402.12348*, 2024.
- 524 Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha
 525 Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. The llama 3 herd of models.
 526 *arXiv preprint arXiv:2407.21783*, 2024.
 527
- 528 Izzeddin Gur, Hiroki Furuta, Austin Huang, Mustafa Safdari, Yutaka Matsuo, Douglas Eck, and
 529 Aleksandra Faust. A real-world webagent with planning, long context understanding, and pro-
 530 gram synthesis. *arXiv preprint arXiv:2307.12856*, 2023.
- 531 Dongge Han, Trevor McInroe, Adam Jolley, Stefano V Albrecht, Peter Bell, and Amos Storkey.
 532 Llm-personalize: Aligning llm planners with human preferences via reinforced self-training for
 533 housekeeping robots. *arXiv preprint arXiv:2404.14285*, 2024.
- 534 Shibo Hao, Yi Gu, Haodi Ma, Joshua Jiahua Hong, Zhen Wang, Daisy Zhe Wang, and Zhiting Hu.
 535 Reasoning with language model is planning with world model. *arXiv preprint arXiv:2305.14992*,
 536 2023.
 537
- 538 Matthew Hausknecht, Prithviraj Ammanabrolu, Marc-Alexandre Côté, and Xingdi Yuan. Interac-
 539 tive fiction games: A colossal adventure. In *Proceedings of the AAAI Conference on Artificial
 Intelligence*, volume 34, pp. 7903–7910, 2020.

- 540 Wenyue Hua, Lizhou Fan, Lingyao Li, Kai Mei, Jianchao Ji, Yingqiang Ge, Libby Hemphill, and
 541 Yongfeng Zhang. War and peace (waragent): Large language model-based multi-agent simulation
 542 of world wars. *arXiv preprint arXiv:2311.17227*, 2023.
- 543
- 544 Xu Huang, Weiwen Liu, Xiaolong Chen, Xingmei Wang, Hao Wang, Defu Lian, Yasheng Wang,
 545 Ruiming Tang, and Enhong Chen. Understanding the planning of llm agents: A survey. *arXiv
 546 preprint arXiv:2402.02716*, 2024a.
- 547 Yuzhen Huang, Yuzhuo Bai, Zhihao Zhu, Junlei Zhang, Jinghan Zhang, Tangjun Su, Junteng Liu,
 548 Chuancheng Lv, Yikai Zhang, Yao Fu, et al. C-eval: A multi-level multi-discipline chinese eval-
 549 uation suite for foundation models. *Advances in Neural Information Processing Systems*, 36,
 550 2024b.
- 551 Ziqi Jin and Wei Lu. Tab-cot: Zero-shot tabular chain of thought. *arXiv preprint arXiv:2305.17812*,
 552 2023.
- 553
- 554 Hanyu Lai, Xiao Liu, Iat Long Iong, Shuntian Yao, Yuxuan Chen, Pengbo Shen, Hao Yu, Hanchen
 555 Zhang, Xiaohan Zhang, Yuxiao Dong, and Jie Tang. Autowebglm: Bootstrap and reinforce a
 556 large language model-based web navigating agent, 2024. URL <https://arxiv.org/abs/2404.03648>.
- 557
- 558 Minghao Li, Yingxiu Zhao, Bowen Yu, Feifan Song, Hangyu Li, Haiyang Yu, Zhoujun Li, Fei
 559 Huang, and Yongbin Li. Api-bank: A comprehensive benchmark for tool-augmented llms, 2023a.
 560 URL <https://arxiv.org/abs/2304.08244>.
- 561
- 562 Wenhao Li, Zhiyuan Yu, Qijin She, Zhinan Yu, Yuqing Lan, Chenyang Zhu, Ruizhen Hu, and Kai
 563 Xu. Llm-enhanced scene graph learning for household rearrangement, 2024. URL <https://arxiv.org/abs/2408.12093>.
- 564
- 565 Yang Li, Yangyang Yu, Haohang Li, Zhi Chen, and Khaldoun Khashanah. Tradinggpt: Multi-agent
 566 system with layered memory and distinct characters for enhanced financial trading performance.
 567 *arXiv preprint arXiv:2309.03736*, 2023b.
- 568
- 569 Hao Liu, Lisa Lee, Kimin Lee, and Pieter Abbeel. Instruction-following agents with multimodal
 570 transformer. *arXiv preprint arXiv:2210.13431*, 2022.
- 571
- 572 Shaoteng Liu, Haoqi Yuan, Minda Hu, Yanwei Li, Yukang Chen, Shu Liu, Zongqing Lu, and
 573 Jiaya Jia. Rl-gpt: Integrating reinforcement learning and code-as-policy. *arXiv preprint
 574 arXiv:2402.19299*, 2024.
- 575
- 576 Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding,
 577 Kaiwen Men, Kejuan Yang, et al. Agentbench: Evaluating llms as agents. *arXiv preprint
 578 arXiv:2308.03688*, 2023a.
- 579
- 580 Zhiwei Liu, Weiran Yao, Jianguo Zhang, Le Xue, Shelby Heinecke, Rithesh Murthy, Yihao Feng,
 581 Zeyuan Chen, Juan Carlos Niebles, Devansh Arpit, et al. Bolaa: Benchmarking and orchestrating
 582 llm-augmented autonomous agents. *arXiv preprint arXiv:2308.05960*, 2023b.
- 583
- 584 Albert Lu, Hongxin Zhang, Yanzhe Zhang, Xuezhi Wang, and Diyi Yang. Bounding the capabili-
 585 ties of large language models in open text generation with prompt constraints. *arXiv preprint
 586 arXiv:2302.09185*, 2023.
- 587
- 588 Chang Ma, Junlei Zhang, Zhihao Zhu, Cheng Yang, Yujiu Yang, Yaohui Jin, Zhenzhong Lan, Ling-
 589 peng Kong, and Junxian He. Agentboard: An analytical evaluation board of multi-turn llm agents.
 590 *arXiv preprint arXiv:2401.13178*, 2024.
- 591
- 592 Weiyu Ma, Qirui Mi, Xue Yan, Yuqiao Wu, Runji Lin, Haifeng Zhang, and Jun Wang. Large
 593 language models play starcraft ii: Benchmarks and a chain of summarization approach. *arXiv
 594 preprint arXiv:2312.11865*, 2023.
- 595
- 596 Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri
 597 Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. Self-refine: Iterative refinement
 598 with self-feedback. *Advances in Neural Information Processing Systems*, 36, 2024.

- 594 Meredith Ringel Morris, Jascha Sohl-Dickstein, Noah Fiedel, Tris Warkentin, Allan Dafoe, Alek-
 595 sandra Faust, Clement Farabet, and Shane Legg. Position: Levels of agi for operationalizing
 596 progress on the path to agi. In *Forty-first International Conference on Machine Learning*.
 597
- 598 Gabriel Mukobi, Hannah Erlebach, Niklas Lauffer, Lewis Hammond, Alan Chan, and Jesse
 599 Clifton. Welfare diplomacy: Benchmarking language model cooperation. *arXiv preprint*
 600 *arXiv:2310.08901*, 2023.
- 601 OpenAI. Gpt-3.5 turbo. <https://platform.openai.com/docs/models/gpt-3-5>,
 602 2023.
- 603 OpenAI. Openai o1-mini: Advancing cost-efficient reasoning, 2024a. URL <https://openai.com/index/openai-o1-mini-advancing-cost-efficient-reasoning/>.
- 604 OpenAI. Hello gpt-4o, 2024b. URL <https://openai.com/index/hello-gpt-4o>.
- 605 Aloïs Pourchot and Olivier Sigaud. Cem-rl: Combining evolutionary and gradient-based methods
 606 for policy search. *arXiv preprint arXiv:1810.01222*, 2018.
- 607 David Premack and Guy Woodruff. Does the chimpanzee have a theory of mind? *Behavioral and*
 608 *brain sciences*, 1(4):515–526, 1978.
- 609 Siyuan Qi, Shuo Chen, Yexin Li, Xiangyu Kong, Junqi Wang, Bangcheng Yang, Pring Wong, Yifan
 610 Zhong, Xiaoyuan Zhang, Zhaowei Zhang, et al. Civrealm: A learning and reasoning odyssey in
 611 civilization for decision-making agents. *arXiv preprint arXiv:2401.10568*, 2024.
- 612 Chen Qian, Wei Liu, Hongzhang Liu, Nuo Chen, Yufan Dang, Jiahao Li, Cheng Yang, Weize Chen,
 613 Yusheng Su, Xin Cong, et al. Chatdev: Communicative agents for software development. In *Pro-
 614 ceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume
 615 1: Long Papers)*, pp. 15174–15186, 2024.
- 616 Machel Reid, Nikolay Savinov, Denis Teplyashin, Dmitry Lepikhin, Timothy Lillicrap, Jean-
 617 baptiste Alayrac, Radu Soricut, Angeliki Lazaridou, Orhan Firat, Julian Schrittweis, et al. Gemini
 618 1.5: Unlocking multimodal understanding across millions of tokens of context. *arXiv preprint*
 619 *arXiv:2403.05530*, 2024.
- 620 Mikayel Samvelyan, Tabish Rashid, Christian Schroeder De Witt, Gregory Farquhar, Nantas
 621 Nardelli, Tim GJ Rudner, Chia-Man Hung, Philip HS Torr, Jakob Foerster, and Shimon Whiteson.
 622 The starcraft multi-agent challenge. *arXiv preprint arXiv:1902.04043*, 2019.
- 623 Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion:
 624 Language agents with verbal reinforcement learning. *Advances in Neural Information Processing
 625 Systems*, 36, 2024.
- 626 Aarohi Srivastava, Abhinav Rastogi, Abhishek Rao, Abu Awal Md Shoeb, Abubakar Abid, Adam
 627 Fisch, Adam R Brown, Adam Santoro, Aditya Gupta, Adrià Garriga-Alonso, et al. Beyond the
 628 imitation game: Quantifying and extrapolating the capabilities of language models. *arXiv preprint*
 629 *arXiv:2206.04615*, 2022.
- 630 Aarohi Srivastava, Abhinav Rastogi, Abhishek Rao, Abu Awal Md Shoeb, Abubakar Abid, Adam
 631 Fisch, Adam R. Brown, Adam Santoro, Aditya Gupta, Adrià Garriga-Alonso, Agnieszka Kluska,
 632 Aitor Lewkowycz, Akshat Agarwal, Alethea Power, Alex Ray, Alex Warstadt, Alexander W.
 633 Kocurek, Ali Safaya, Ali Tazarv, Alice Xiang, Alicia Parrish, Allen Nie, Aman Hussain,
 634 Amanda Askell, Amanda Dsouza, Ambrose Slone, Ameet Rahane, Anantharaman S. Iyer, An-
 635 ders Andreassen, Andrea Madotto, Andrea Santilli, Andreas Stuhlmüller, Andrew Dai, An-
 636 drew La, Andrew Lampinen, Andy Zou, Angela Jiang, Angelica Chen, Anh Vuong, Animesh
 637 Gupta, Anna Gottardi, Antonio Norelli, Anu Venkatesh, Arash Gholamidavoodi, Arfa Tabas-
 638 sum, Arul Menezes, Arun Kirubarajan, Asher Mollokandov, Ashish Sabharwal, Austin Her-
 639 rick, Avia Efrat, Aykut Erdem, Ayla Karakaş, B. Ryan Roberts, Bao Sheng Loe, Barret Zoph,
 640 Bartłomiej Bojanowski, Batuhan Özyurt, Behnam Hedayatnia, Behnam Neyshabur, Benjamin
 641 Inden, Benno Stein, Berk Ekmekci, Bill Yuchen Lin, Blake Howald, Bryan Orinon, Cameron
 642 Diao, Cameron Dour, Catherine Stinson, Cedrick Argueta, César Ferri Ramírez, Chandan Singh,
 643 Charles Rathkopf, Chenlin Meng, Chitta Baral, Chiyu Wu, Chris Callison-Burch, Chris Waites,

648 Christian Voigt, Christopher D. Manning, Christopher Potts, Cindy Ramirez, Clara E. Rivera,
 649 Clemencia Siro, Colin Raffel, Courtney Ashcraft, Cristina Garbacea, Damien Sileo, Dan Gar-
 650 rette, Dan Hendrycks, Dan Kilman, Dan Roth, Daniel Freeman, Daniel Khashabi, Daniel Levy,
 651 Daniel Moseguí González, Danielle Perszyk, Danny Hernandez, Danqi Chen, Daphne Ippolito,
 652 Dar Gilboa, David Dohan, David Drakard, David Jurgens, Debajyoti Datta, Deep Ganguli, De-
 653 nis Emelin, Denis Kleyko, Deniz Yuret, Derek Chen, Derek Tam, Dieuwke Hupkes, Diganta
 654 Misra, Dilyar Buzan, Dimitri Coelho Mollo, Diyi Yang, Dong-Ho Lee, Dylan Schrader, Eka-
 655 terina Shutova, Ekin Dogus Cubuk, Elad Segal, Eleanor Hagerman, Elizabeth Barnes, Eliza-
 656 beth Donoway, Ellie Pavlick, Emanuele Rodola, Emma Lam, Eric Chu, Eric Tang, Erkut Erdem,
 657 Ernie Chang, Ethan A. Chi, Ethan Dyer, Ethan Jerzak, Ethan Kim, Eunice Engefu Manyasi, Ev-
 658 genii Zheltonozhskii, Fanyue Xia, Fatemeh Siar, Fernando Martínez-Plumed, Francesca Happé,
 659 Francois Chollet, Frieda Rong, Gaurav Mishra, Genta Indra Winata, Gerard de Melo, Germán
 660 Kruszewski, Giambattista Parascandolo, Giorgio Mariani, Gloria Wang, Gonzalo Jaimovich-
 661 López, Gregor Betz, Guy Gur-Ari, Hana Galijasevic, Hannah Kim, Hannah Rashkin, Hannaneh
 662 Hajishirzi, Harsh Mehta, Hayden Bogar, Henry Shevlin, Hinrich Schütze, Hiromu Yakura, Hong-
 663 ming Zhang, Hugh Mee Wong, Ian Ng, Isaac Noble, Jaap Jumelet, Jack Geissinger, Jackson
 664 Kernion, Jacob Hilton, Jaehoon Lee, Jaime Fernández Fisac, James B. Simon, James Koppel,
 665 James Zheng, James Zou, Jan Kocoń, Jana Thompson, Janelle Wingfield, Jared Kaplan, Jarema
 666 Radom, Jascha Sohl-Dickstein, Jason Phang, Jason Wei, Jason Yosinski, Jekaterina Novikova,
 667 Jelle Bosscher, Jennifer Marsh, Jeremy Kim, Jeroen Taal, Jesse Engel, Jesujoba Alabi, Ji-
 668 acheng Xu, Jiaming Song, Jillian Tang, Joan Waweru, John Burden, John Miller, John U. Balis,
 669 Jonathan Batchelder, Jonathan Berant, Jörg Frohberg, Jos Rozen, Jose Hernandez-Orallo, Joseph
 670 Boudeman, Joseph Guerr, Joseph Jones, Joshua B. Tenenbaum, Joshua S. Rule, Joyce Chua,
 671 Kamil Kanclerz, Karen Livescu, Karl Krauth, Karthik Gopalakrishnan, Katerina Ignatyeva, Katja
 672 Markert, Kaustubh D. Dhole, Kevin Gimpel, Kevin Omondi, Kory Mathewson, Kristen Chia-
 673 fullo, Ksenia Shkaruta, Kumar Shridhar, Kyle McDonell, Kyle Richardson, Laria Reynolds, Leo
 674 Gao, Li Zhang, Liam Dugan, Lianhui Qin, Lidia Contreras-Ochando, Louis-Philippe Morency,
 675 Luca Moschella, Lucas Lam, Lucy Noble, Ludwig Schmidt, Luheng He, Luis Oliveros Colón,
 676 Luke Metz, Lütfi Kerem Şenel, Maarten Bosma, Maarten Sap, Maartje ter Hoeve, Maheen Fa-
 677 rooqi, Manaal Faruqui, Mantas Mazeika, Marco Baturan, Marco Marelli, Marco Maru, Maria
 678 Jose Ramírez Quintana, Marie Tolkiehn, Mario Julianelli, Martha Lewis, Martin Potthast,
 679 Matthew L. Leavitt, Matthias Hagen, Mátyás Schubert, Medina Orduna Baitemirova, Melody
 680 Arnaud, Melvin McElrath, Michael A. Yee, Michael Cohen, Michael Gu, Michael Ivanitskiy,
 681 Michael Starritt, Michael Strube, Michał Swedrowski, Michele Bevilacqua, Michihiro Yasunaga,
 682 Mihir Kale, Mike Cain, Mimeo Xu, Mirac Suzgun, Mitch Walker, Mo Tiwari, Mohit Bansal,
 683 Moin Aminnaseri, Mor Geva, Mozhdeh Gheini, Mukund Varma T, Nanyun Peng, Nathan A.
 684 Chi, Nayeon Lee, Neta Gur-Ari Krakover, Nicholas Cameron, Nicholas Roberts, Nick Doiron,
 685 Nicole Martinez, Nikita Nangia, Niklas Deckers, Niklas Muennighoff, Nitish Shirish Keskar,
 686 Niveditha S. Iyer, Noah Constant, Noah Fiedel, Nuan Wen, Oliver Zhang, Omar Agha, Omar El-
 687 baghdadi, Omer Levy, Owain Evans, Pablo Antonio Moreno Casares, Parth Doshi, Pascale Fung,
 688 Paul Pu Liang, Paul Vicol, Pegah Alipoormolabashi, Peiyuan Liao, Percy Liang, Peter Chang, Pe-
 689 ter Eckersley, Phu Mon Htut, Pinyu Hwang, Piotr Miłkowski, Piyush Patil, Pouya Pezeshkpour,
 690 Priti Oli, Qiaozhu Mei, Qing Lyu, Qinlang Chen, Rabin Banjade, Rachel Etta Rudolph, Raefer
 691 Gabriel, Rahel Habacker, Ramon Risco, Raphaël Millière, Rhythm Garg, Richard Barnes, Rif A.
 692 Saurous, Riku Arakawa, Robbe Raymaekers, Robert Frank, Rohan Sikand, Roman Novak, Ro-
 693 man Sitelew, Ronan LeBras, Rosanne Liu, Rowan Jacobs, Rui Zhang, Ruslan Salakhutdinov,
 694 Ryan Chi, Ryan Lee, Ryan Stovall, Ryan Teehan, Rylan Yang, Sahib Singh, Saif M. Moham-
 695 mad, Sajant Anand, Sam Dillavou, Sam Shleifer, Sam Wiseman, Samuel Gruetter, Samuel R.
 696 Bowman, Samuel S. Schoenholz, Sanghyun Han, Sanjeev Kwatra, Sarah A. Rous, Sarik Ghaz-
 697 arian, Sayan Ghosh, Sean Casey, Sebastian Bischoff, Sebastian Gehrmann, Sebastian Schus-
 698 ter, Sepideh Sadeghi, Shadi Hamdan, Sharon Zhou, Shashank Srivastava, Sherry Shi, Shikhar
 699 Singh, Shima Asaadi, Shixiang Shane Gu, Shubh Pachchigar, Shubham Toshniwal, Shyam Upad-
 700 hyay, Shyamolima, Debnath, Siamak Shakeri, Simon Thormeyer, Simone Melzi, Siva Reddy,
 701 Sneha Priscilla Makini, Soo-Hwan Lee, Spencer Torene, Sriharsha Hatwar, Stanislas Dehaene,
 Stefan Divic, Stefano Ermon, Stella Biderman, Stephanie Lin, Stephen Prasad, Steven T. Pianta-
 dosi, Stuart M. Shieber, Summer Misherghi, Svetlana Kiritchenko, Swaroop Mishra, Tal Linzen,
 Tal Schuster, Tao Li, Tao Yu, Tariq Ali, Tatsu Hashimoto, Te-Lin Wu, Théo Desbordes, Theodore
 Rothschild, Thomas Phan, Tianle Wang, Tiberius Nkinyili, Timo Schick, Timofei Kornev, Ti-
 tus Tunduny, Tobias Gerstenberg, Trenton Chang, Trishala Neeraj, Tushar Khot, Tyler Shultz,

- 702 Uri Shaham, Vedant Misra, Vera Demberg, Victoria Nyamai, Vikas Raunak, Vinay Ramasesh,
 703 Vinay Uday Prabhu, Vishakh Padmakumar, Vivek Srikumar, William Fedus, William Saun-
 704 ders, William Zhang, Wout Vossen, Xiang Ren, Xiaoyu Tong, Xinran Zhao, Xinyi Wu, Xudong
 705 Shen, Yadollah Yaghoobzadeh, Yair Lakretz, Yangqiu Song, Yasaman Bahri, Yejin Choi, Yichi
 706 Yang, Yiding Hao, Yifu Chen, Yonatan Belinkov, Yu Hou, Yufang Hou, Yuntao Bai, Zachary
 707 Seid, Zhuoye Zhao, Zijian Wang, Zijie J. Wang, Zirui Wang, and Ziyi Wu. Beyond the im-
 708 itation game: Quantifying and extrapolating the capabilities of language models, 2023. URL
 709 <https://arxiv.org/abs/2206.04615>.
- 710 Winnie Street, John Oliver Siy, Geoff Keeling, Adrien Baranes, Benjamin Barnett, Michael McK-
 711 ibben, Tatenda Kanyere, Alison Lentz, Robin IM Dunbar, et al. Llms achieve adult human per-
 712 formance on higher-order theory of mind tasks. *arXiv preprint arXiv:2405.18870*, 2024.
- 713 Qiaoyu Tang, Ziliang Deng, Hongyu Lin, Xianpei Han, Qiao Liang, Boxi Cao, and Le Sun. Toolal-
 714 paca: Generalized tool learning for language models with 3000 simulated cases. *arXiv preprint*
 715 *arXiv:2306.05301*, 2023.
- 716 Oriol Vinyals, Igor Babuschkin, Wojciech M Czarnecki, Michaël Mathieu, Andrew Dudzik, Juny-
 717 oung Chung, David H Choi, Richard Powell, Timo Ewalds, Petko Georgiev, et al. Grandmaster
 718 level in starcraft ii using multi-agent reinforcement learning. *nature*, 575(7782):350–354, 2019.
- 719 Alex Wang, Yada Pruksachatkun, Nikita Nangia, Amanpreet Singh, Julian Michael, Felix Hill, Omer
 720 Levy, and Samuel Bowman. Super glue: A stickier benchmark for general-purpose language
 721 understanding systems. *Advances in neural information processing systems*, 32, 2019.
- 722 Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan,
 723 and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models.
 724 *arXiv preprint arXiv:2305.16291*, 2023.
- 725 Ruoyo Wang, Peter Jansen, Marc-Alexandre Côté, and Prithviraj Ammanabrolu. Scienceworld: Is
 726 your agent smarter than a 5th grader? *arXiv preprint arXiv:2203.07540*, 2022.
- 727 Jason Wei, Maarten Bosma, Vincent Y Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du,
 728 Andrew M Dai, and Quoc V Le. Finetuned language models are zero-shot learners. *arXiv preprint*
 729 *arXiv:2109.01652*, 2021.
- 730 Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny
 731 Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in*
 732 *neural information processing systems*, 35:24824–24837, 2022.
- 733 Martin Weiss, Nasim Rahaman, Manuel Wuthrich, Yoshua Bengio, Li Erran Li, Bernhard
 734 Schölkopf, and Christopher Pal. Rethinking the buyer’s inspection paradox in information mar-
 735 kets with language agents.
- 736 Wikipedia. Civilization, 2024a. URL <https://en.wikipedia.org/wiki/Civilization>. Accessed: 2024-09-23.
- 737 Wikipedia. Diplomacy, 2024b. URL [https://en.wikipedia.org/wiki/Diplomacy_\(game\)](https://en.wikipedia.org/wiki/Diplomacy_(game)). Accessed: 2024-09-23.
- 738 Wikipedia. Honor of kings, 2024c. URL https://en.wikipedia.org/wiki/Honor_of_Kings. Accessed: 2024-09-23.
- 739 Wikipedia. Mafia (party game), 2024d. URL [https://en.wikipedia.org/wiki/Mafia_\(party_game\)](https://en.wikipedia.org/wiki/Mafia_(party_game)). Accessed: 2024-09-23.
- 740 Wikipedia. Starcraft ii, 2024e. URL https://en.wikipedia.org/wiki/StarCraft_II. Accessed: 2024-09-23.
- 741 Wikipedia. Dota 2, 2024f. URL https://en.wikipedia.org/wiki/Dota_2. Accessed:
 742 2024-09-23.
- 743 Michael Wooldridge and Nicholas R Jennings. Intelligent agents: Theory and practice. *The knowl-
 744 edge engineering review*, 10(2):115–152, 1995.

- 756 Yue Wu, Xuan Tang, Tom M Mitchell, and Yuanzhi Li. Smartplay: A benchmark for llms as
 757 intelligent agents. *arXiv preprint arXiv:2310.01557*, 2023.
- 758
- 759 Zhiheng Xi, Wenxiang Chen, Xin Guo, Wei He, Yiwen Ding, Boyang Hong, Ming Zhang, Junzhe
 760 Wang, Senjie Jin, Enyu Zhou, et al. The rise and potential of large language model based agents:
 761 A survey. *arXiv preprint arXiv:2309.07864*, 2023.
- 762 Zhiheng Xi, Yiwen Ding, Wenxiang Chen, Boyang Hong, Honglin Guo, Junzhe Wang, Dingwen
 763 Yang, Chenyang Liao, Xin Guo, Wei He, et al. Agentgym: Evolving large language model-based
 764 agents across diverse environments. *arXiv preprint arXiv:2406.04151*, 2024.
- 765
- 766 Bushi Xiao, Ziyuan Yin, and Zixuan Shan. Simulating public administration crisis: A novel genera-
 767 tive agent-based simulation system to lower technology barriers in social science research. *arXiv*
 768 *preprint arXiv:2311.06957*, 2023.
- 769 Jian Xie, Kai Zhang, Jiangjie Chen, Tinghui Zhu, Renze Lou, Yuandong Tian, Yanghua Xiao, and
 770 Yu Su. Travelplanner: A benchmark for real-world planning with language agents. *arXiv preprint*
 771 *arXiv:2402.01622*, 2024.
- 772 Lin Xu, Zhiyuan Hu, Daquan Zhou, Hongyu Ren, Zhen Dong, Kurt Keutzer, See-Kiong Ng, and
 773 Jiaoshi Feng. Magic: Investigation of large language model powered multi-agent in cognition,
 774 adaptability, rationality and collaboration. In *ICLR 2024 Workshop on Large Language Model*
 775 (*LLM*) *Agents*, 2023.
- 776
- 777 Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao.
 778 React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*,
 779 2022.
- 780 Shunyu Yao, Howard Chen, John Yang, and Karthik Narasimhan. Webshop: Towards scalable
 781 real-world web interaction with grounded language agents. In *ArXiv*, preprint.
- 782
- 783 Jiaxuan You, Ge Liu, Yunzhu Li, Song Han, and Dawn Song. How far are we from agi. In *ICLR*
 784 *2024 Workshops*, 2024.
- 785 Xiang Yue, Yuansheng Ni, Kai Zhang, Tianyu Zheng, Ruqi Liu, Ge Zhang, Samuel Stevens,
 786 Dongfu Jiang, Weiming Ren, Yuxuan Sun, et al. Mmmu: A massive multi-discipline multi-
 787 modal understanding and reasoning benchmark for expert agi. In *Proceedings of the IEEE/CVF*
 788 *Conference on Computer Vision and Pattern Recognition*, pp. 9556–9567, 2024.
- 789
- 790 Wenqi Zhang, Ke Tang, Hai Wu, Mengna Wang, Yongliang Shen, Guiyang Hou, Zeqi Tan, Peng Li,
 791 Yueting Zhuang, and Weiming Lu. Agent-pro: Learning to evolve via policy-level reflection and
 792 optimization. *arXiv preprint arXiv:2402.17574*, 2024a.
- 793 Zeyu Zhang, Xiaohe Bo, Chen Ma, Rui Li, Xu Chen, Quanyu Dai, Jieming Zhu, Zhenhua Dong,
 794 and Ji-Rong Wen. A survey on the memory mechanism of large language model based agents.
 795 *arXiv preprint arXiv:2404.13501*, 2024b.
- 796
- 797 Zhuosheng Zhang, Yao Yao, Aston Zhang, Xiangru Tang, Xinbei Ma, Zhiwei He, Yiming Wang,
 798 Mark Gerstein, Rui Wang, Gongshen Liu, et al. Igniting language intelligence: The hitchhiker’s
 799 guide from chain-of-thought reasoning to language agents. *arXiv preprint arXiv:2311.11797*,
 2023.
- 800
- 801 Longtao Zheng, Rundong Wang, Xinrun Wang, and Bo An. Synapse: Trajectory-as-exemplar
 802 prompting with memory for computer control. In *The Twelfth International Conference on Learn-
 803 ing Representations*, 2023a.
- 804
- 805 Qinkai Zheng, Xiao Xia, Xu Zou, Yuxiao Dong, Shan Wang, Yufei Xue, Zihan Wang, Lei Shen,
 806 Andi Wang, Yang Li, et al. Codegeex: A pre-trained model for code generation with multilingual
 807 evaluations on humaneval-x. *arXiv preprint arXiv:2303.17568*, 2023b.
- 808
- 809

810 A APPENDIX
811812 A.1 EVALUATION METRICS
813

	Capacity	Computer Formula
Strategic planning	$RPM_i = \sum_{t=1}^T (\text{collected_minerals}_i(t) + \text{collected_vespene}_i(t))$	
	$EER_i = \frac{\sum_{t=1}^T (\text{collected_minerals}_i(t) + \text{spent_vespene}_i(t)) \times 100}{\sum_{t=1}^T (\text{collected_minerals}_i(t) + \text{collected_vespene}_i(t))}$	
	$SUR_i = \frac{\sum_{t=1}^T \text{supply_used}_i(t)}{\sum_{t=1}^T \text{supply_cap}_i(t)}$	
	$TRR_i = \frac{\text{completed_tech}_i}{\text{total_research_count}_i}$	
Real-time decision making	$APM_i = \frac{\text{total_actions}_i}{\text{game_time_minutes}_i}$	
	$EPM_i = \frac{\text{effective_actions}_i}{\text{game_time_minutes}_i}$	
Adaptability	$\text{WinRateTrend}_i = \frac{\text{WinRate}_{\text{end},i} - \text{WinRate}_{\text{start},i}}{\text{total_games}}$	
	$\text{ErrorRateTrend}_i = \frac{\text{ErrorRate}_{\text{start},i} - \text{ErrorRate}_{\text{end},i}}{\text{total_games}}$	

814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
Table 5: This table presents the key capacity metrics and their corresponding computational formulas used to evaluate LLMs in StarCraft II. The metrics are categorized under three primary capacities: Strategic Planning, Real-time Decision Making, and Adaptability. Each metric captures different aspects of the LLM’s performance, such as resource management (RPM), supply utilization (SUR), action efficiency (APM, EPM), and adaptation trends (WinRateTrend, ErrorRateTrend), providing a comprehensive assessment of the model’s gameplay capabilities.

A.2 GAMES INTRODUCTION

StarCraft II is a real-time strategy game whose core mechanics include resource management, base building, troop production and command. Players need to efficiently gather resources, build and upgrade bases, train various military units, and defeat opponents through precise micromanagement and macro-strategy in a real-time environment. The game emphasizes quick decision-making and flexibility, requiring players to balance economic development and military operations in a highly dynamic battlefield in order to ultimately destroy the enemy’s base and win.

The core mechanics of Civilization revolve around turn-based strategy, where players lead a civilization from antiquity to the future by managing cities, developing technology and culture, exploring maps, and engaging in diplomacy and warfare. With an emphasis on resource management, long-term planning, and strategic decision-making, the game requires players to unlock new abilities through the tech and culture trees, choose different victory conditions (e.g., military victory, tech victory, or cultural victory), and gain an advantage in their interactions with other civilizations. The variety and depth of the game makes it a classic strategy game.

864
865
866
867
868
869

Game manual		
Setting	Map specification	Standard 1v1 map with mining, gas, expansion points, obstacle terrain and other elements (e.g. map: Jagannatha LE).
	Number of players	2 players per match against each other.
	Resource type	Two main resources - minerals and gases, used for unit production and technological upgrading.
Unit configuration and policy	Basic unit configuration	12 farmers (SCV/Probe/Drone) for resource collection. 1 main base (Command Center/Nexus/Hatchery). 1 Supply Depot (Pylon/Overlord) to control the population cap.
		Terran: Focuses on mechanical units and air power, with strong defensive and multi-functional building capabilities.
	Ethnic divisions	Protoss: has shields and powerful individual units, but is slower to produce. Zerg: Unit production is fast, relying on massive ground forces and good ecological control.
	Unit Production and Technology tree	Terran: can produce ground units (such as Marine, Marauder) and air force units (such as Viking, Banshee). Protoss: Can produce high-attack units (e.g., Zealot, Stalker) and powerful air units (e.g., Carrier, Phoenix). Zerg: Can produce a large number of cheap units (such as Zergling, Hydralisk) and high-tech units (such as Mutalisk, Ultralisk).
Fixed opening strategy	Initial base strategy	Rapid expansion strategy: quickly establish a second base to enhance economic output and increase resource collection speed. Quick attack strategy: Quickly produce early combat units, directly attack enemy bases, forcing opponents to defend.
		Defensive strategy: Strengthen fortifications (such as Terran's Bunker, Protoss 'Photon Cannon) to delay enemy attacks and save strength for later development.
	Army layout and defense	Defensive arrangement: Arrange defensive units near the base to ensure the safety of the mining area and the main base; Different races have different defensive structures, such as Terran's Bunker, Protoss 'Shield Battery, and Zerg's Spine Crawler. Offensive placement: Deploy units to harass and control key locations on the map, such as enemy resource points.
Goals	Economic development	Through the collection of minerals and gases, the rapid development of the economy and science and technology, the establishment of a more powerful army.
	Military victory	Destroy all their main bases or render them incapable of reproducing units.
	Map control	Capture key positions on the map (e.g., resource points, highlands) and use tactical advantage to overwhelm the opponent's economy.

911 Table 6: Game manual detailing the settings, unit configuration and policy, fixed opening strategies, and goals for gameplay in StarCraft II. This table provides a comprehensive overview of the
912 game mechanics, including map specifications, player setup, unit production capabilities, strategic
913 approaches, and objectives essential for effective gameplay.
914

915
916
917